

Team:

Khantil Choksi - khchoksi	Shubhankar Reddy - skatta2
---------------------------	----------------------------

Staging Server IP			VM IP		
152.14.83.156	ece792	EcE792net!	192.168.124.15	ece792	EcE792net!

PROBLEM 1:

	500 Mbps	1000 Mbps	5000 Mbps
CPU	1 vCPU	1 vCPU (2 for the AX tech stack)	1/2/8 vCPU (for IP base, Security and AppX)
Memory	4 GB	4 GB	4 GB

- a) Networking: BGP, OSPF, EIGRP, Routing Information Protocol (RIP), Intermediate System-to-Intermediate System (IS-IS), IPv6, GRE, VRF-Lite, NTP
- b) Security: ZBFW, IPsec VPN, Easy VPN, DMVPN, FlexVPN
- c) Management: Cisco IOS XE CLI, SSH, Flexible NetFlow, SNMP, EEM, and NETCONF

1 year annual cost of CSR 1000V (AX stack) on aws = \$ 3,723.00 (Software cost)

1 year of c4.large EC2 instance cost = \$515.00 (C4.large, yearly cost)

Cost of running two 1000 Mbps CSR in AWS = $(3,723 + 515.00) * 2 = \$8,476$ USD

https://aws.amazon.com/marketplace/pp/B00OCG4OAA?qid=1539727240910&sr=0-1&ref_=srh_res_product_title&cl_spe=T1

<https://aws.amazon.com/ec2/pricing/reserved-instances/pricing/>

PROBLEM 2:

```
sudo virt-install -n khchoksi -r 2048 --vcpu=4 --cpu host --disk  
path=/var/lib/libvirt/images/khchoksi.img,size=10 --network network=khchoksivm -c  
/home/ece792/iso/CentOS-7-x86_64-Minimal-1804.iso -v
```

```
ece792@ece792-Standard-PC-i440FX-PIIX-1996:~$ virsh list --all  
Id      Name                               State  
-----  
3       khchoksi                           running  
  
ece792@ece792-Standard-PC-i440FX-PIIX-1996:~$ virsh domifaddr khchoksi  
Name      MAC address      Protocol  Address  
-----  
vnet0     52:54:00:2e:29:b6  ipv4      192.168.122.92/24  
  
ece792@ece792-Standard-PC-i440FX-PIIX-1996:~$
```

```
CentOS Linux 7 (Core)  
Kernel 3.10.0-862.el7.x86_64 on an x86_64  
  
localhost login:  
  
CentOS Linux 7 (Core)  
Kernel 3.10.0-862.el7.x86_64 on an x86_64  
  
localhost login:  
root  
  
CentOS Linux 7 (Core)  
Kernel 3.10.0-862.el7.x86_64 on an x86_64  
  
localhost login: Password:  
Login incorrect  
  
localhost login: root  
Password:  
Last failed login: Sat Oct 6 15:49:48 EDT 2018 on tty1  
There was 1 failed login attempt since the last successful login.  
[root@localhost ~]#  
[root@localhost ~]#  
[root@localhost ~]#  
[root@localhost ~]#  
[root@localhost ~]#  
[root@localhost ~]#  
[root@localhost ~]# dhclient  
  
[root@localhost ~]#  
[root@localhost ~]# _
```

To install required applications:

```
$ yum install iperf3
```

```
$ yum install wireshark
```

(i)

	VM's NIC	hypervisor NIC
IP Address	192.168.122.92/24	192.168.124.15
MAC Address	52:54:00:2e:29:b6	52:54:00:2f:dd:ba

```
Complete!
[root@localhost ~]# wireshark
-bash: wireshark: command not found
[root@localhost ~]# tshark
Running as user "root" and group "root". This could be dangerous.
Capturing on 'eth0'
 1 0.000000000 fe:54:00:2e:29:b6 -> Spanning-tree-(for-bridges)_00 STP 52 Conf . Root = 32768/0/fe:54:00:2e:29:b6 Cost = 0 Port = 0x0001
 2 1.984065672 fe:54:00:2e:29:b6 -> Spanning-tree-(for-bridges)_00 STP 52 Conf . Root = 32768/0/fe:54:00:2e:29:b6 Cost = 0 Port = 0x0001
 3 3.999994360 fe:54:00:2e:29:b6 -> Spanning-tree-(for-bridges)_00 STP 52 Conf . Root = 32768/0/fe:54:00:2e:29:b6 Cost = 0 Port = 0x0001
 4 5.984057324 fe:54:00:2e:29:b6 -> Spanning-tree-(for-bridges)_00 STP 52 Conf . Root = 32768/0/fe:54:00:2e:29:b6 Cost = 0 Port = 0x0001
 5 8.000002192 fe:54:00:2e:29:b6 -> Spanning-tree-(for-bridges)_00 STP 52 Conf . Root = 32768/0/fe:54:00:2e:29:b6 Cost = 0 Port = 0x0001
```

(ii)

	output interface of VM
srcIP	192.168.122.92
destIP	172.217.8.14
srcMAC	52:54:00:2e:29:b6
destMAC	fe:54:00:2e:29:b6

	output interface of hypervisor
srcIP	192.168.122.15
destIP	172.217.8.14
srcMAC	52:54:00:2f:dd:aa
destMAC	52:54:00:3b:8a:fe

- Packet going out of the VM

```
$ sudo tshark -i eth0 -T fields -e ip.src -e eth.src -e ip.dst -e eth.dst -e col.Protocol
```

```
[root@localhost ~]# sudo tshark -i eth0 -T fields -e ip.src -e eth.src -e ip.dst -e eth.dst -e col.Protocol
Running as user "root" and group "root". This could be dangerous.
Capturing on 'eth0'
192.168.122.92 52:54:00:2e:29:b6 172.217.8.14 fe:54:00:2e:29:b6 I
CMP
172.217.8.14 fe:54:00:2e:29:b6 192.168.122.92 52:54:00:2e:29:b6 I
CMP
192.168.122.92 52:54:00:2e:29:b6 192.168.122.1 fe:54:00:2e:29:b6 S
SH
```

- Packet going out of the hypervisor using wireshark

The screenshot shows the Wireshark interface with a packet capture on interface 'ens3'. The packet list on the left shows several ICMP packets. The selected packet (No. 102660) is an ICMP Echo request from 192.168.124.15 to 172.217.8.14. The packet details pane on the right shows the following structure:

- Frame 102660: 98 bytes on wire (784 bits), 98 bytes captured (784 bits) on interface 0
- Ethernet II, Src: RealtekU_2f:dd:ba (52:54:00:2f:dd:ba), Dst: RealtekU_3b:8a:fe (52:54:00:3b:8a:fe)
- Internet Protocol Version 4, Src: 192.168.124.15, Dst: 172.217.8.14
- Internet Control Message Protocol

The packet bytes pane shows the raw data of the ICMP Echo request, including the IP header and ICMP header.

These tuples are in different network. According to datapath, the interface ens3 will do NAT and encap, decap operation and changes the src IP to ens3 IP. And it will be sent out the request over ens3.

PROBLEM 3:

1. khchoksiNETWORK2.xml
 - a. `$ virsh net-define khchoksiNETWORK2.xml`
 - b. `$ brctl addbr sw1`
 - c. `virsh net- start khchoksiNETWORK2`

Name	State	Autostart	Persistent
default	active	yes	yes
khchoksiNETWORK2	active	no	yes

2.
 - a. `virsh attach-interface --domain khchoksi --type bridge --source sw1`
 - b. deactivate VM: `$ virsh shutdown khchoksi`
 - c. Add interface to xml file
 - d. `$ virsh define /etc/libvirt/qemu/khchoksi.xml`
 - e. `$ virsh start khchoksi` (restart VM)

```
ece792@ece792-Standard-PC-i440FX-PIIX-1996:~$ virsh domiflist khchoksi
Interface Type      Source      Model      MAC
-----
vnet0     network default     virtio     52:54:00:2e:29:b6
vnet1     bridge  khchoksiNETWORK2 virtio     52:54:00:dd:70:cb
vnet4     bridge  sw1         rtl8139    52:54:00:ab:7f:fb
```

3. `$ virsh suspend khchoksi`
`$ virt clone --original khchoksi --name khchoksilab2VM2 --auto-clone`
`$ virsh resume khchoksi`
`$ virsh start khchoksilab2VM2`

```
ece792@ece792-Standard-PC-i440FX-PIIX-1996:~$ virsh list --all
Id      Name                                State
-----
3       khchoksi                           running
16      khchoksi4                          running
17      khchoksi5                          running
20      khchoksilab2VM2                    running
-       cloned_vm                          shut off
-       khchoksi3                          shut off
```

4. Assign ip to eth1 : ifconfig ens9 10.0.0.1/24

VM	interface	network	MAC Address	IP Addresses
khchoksi	vnet0	default	52:54:00:2e:29:b6	192.168.122.92
	ens9	khchoksiNETWORK2	52:54:00:ab:7f:fb	10.0.0.2
khchoksilab2VM2	vnet2	default	52:54:00:cd:dc:14	192.168.122.51
	ens9	khchoksiNETWORK2	52:54:00:25:f6:4b	10.0.0.3

```
ece792@ece792-Standard-PC-i440FX-PIIX-1996:~$ virsh domiflist khchoksi
Interface Type      Source      Model      MAC
-----
vnet0     network    default     virtio     52:54:00:2e:29:b6
vnet1     bridge     khchoksiNETWORK2 virtio     52:54:00:dd:70:cb
vnet4     bridge     sw1         rtl8139    52:54:00:ab:7f:fb

ece792@ece792-Standard-PC-i440FX-PIIX-1996:~$ virsh domiflist khchoksilab2VM2
Interface Type      Source      Model      MAC
-----
vnet2     network    default     virtio     52:54:00:cd:dc:14
vnet3     bridge     khchoksiNETWORK2 virtio     52:54:00:21:ba:f1
```

5. Tuple fields on 10.0.0.2 are:

```
[root@localhost ~]# tshark -i ens9 -T fields -e ip.src -e eth.src -e ip.dst -e eth.dst
Running as user "root" and group "root". This could be dangerous.
Capturing on 'ens9'
    fe:54:00:25:f6:4b          01:80:c2:00:00:00
    fe:54:00:25:f6:4b          01:80:c2:00:00:00
10.0.0.2      52:54:00:ab:7f:fb          10.0.0.3      52:54:00:25:f6:4b
10.0.0.3      52:54:00:25:f6:4b          10.0.0.2      52:54:00:ab:7f:fb
```

Tuple fields on 10.0.0.3 are:

```
[root@localhost ~]# tshark -i ens9 -T fields -e ip.src -e eth.src -e ip.dst -e eth.dst
Running as user "root" and group "root". This could be dangerous.
Capturing on 'ens9'
    fe:54:00:ab:7f:fb          01:80:c2:00:00:00
    fe:54:00:ab:7f:fb          01:80:c2:00:00:00
10.0.0.2      52:54:00:ab:7f:fb          10.0.0.3      52:54:00:25:f6:4b
10.0.0.3      52:54:00:25:f6:4b          10.0.0.2      52:54:00:ab:7f:fb
```

The tuple fields do not change as the packets are forwarded over L2 and no encaps/decap takes place.

6. On khchoksilab2VM2

```
4: ens9: <BROADCAST,MULTICAST,UP,LOWER_UP> mtu 1500 qdisc pfifo_fa
st state UP group default qlen 1000
    link/ether 52:54:00:25:f6:4b brd ff:ff:ff:ff:ff:ff
    inet 10.0.0.3/24 brd 10.0.0.255 scope global ens9
        valid_lft forever preferred_lft forever
[root@localhost ~]# iperf3 -c 10.0.0.2 -u
Connecting to host 10.0.0.2, port 5201
[ 4] local 10.0.0.3 port 35453 connected to 10.0.0.2 port 5201
[ ID] Interval           Transfer     Bandwidth       Total Datagrams
[ 4]  0.00-1.00    sec    116 KBytes    949 Kbits/sec    82
[ 4]  1.00-2.00    sec    129 KBytes    1.05 Mbits/sec   91
[ 4]  2.00-3.00    sec    127 KBytes    1.04 Mbits/sec   90
[ 4]  3.00-4.00    sec    129 KBytes    1.05 Mbits/sec   91
[ 4]  4.00-5.00    sec    127 KBytes    1.04 Mbits/sec   90
[ 4]  5.00-6.00    sec    129 KBytes    1.05 Mbits/sec   91
[ 4]  6.00-7.00    sec    127 KBytes    1.04 Mbits/sec   90
[ 4]  7.00-8.00    sec    129 KBytes    1.05 Mbits/sec   91
[ 4]  8.00-9.00    sec    127 KBytes    1.04 Mbits/sec   90
[ 4]  9.00-10.00   sec    129 KBytes    1.05 Mbits/sec   91
-----
[ ID] Interval           Transfer     Bandwidth       Jitter        Lost/Total Datagrams
[ 4]  0.00-10.00   sec    1.24 MBytes    1.04 Mbits/sec    0.183 ms      0/897 (0%)
[ 4] Sent 897 datagrams

iperf Done.
```

On khchoksi

```
4: ens9: <BROADCAST,MULTICAST,UP,LOWER_UP> mtu 1500 qdisc pfifo_fa
st state UP group default qlen 1000
    link/ether 52:54:00:ab:7f:fb brd ff:ff:ff:ff:ff:ff
    inet 10.0.0.2/24 brd 10.0.0.255 scope global ens9
        valid_lft forever preferred_lft forever
[root@localhost ~]# iperf3 -s
-----
Server listening on 5201
-----
Accepted connection from 10.0.0.3, port 59600
[ 5] local 10.0.0.2 port 5201 connected to 10.0.0.3 port 35453
[ ID] Interval           Transfer     Bandwidth       Jitter        Lost/Total Datagrams
[ 5]  0.00-1.00    sec    116 KBytes    950 Kbits/sec    0.488 ms      0/82 (0%)
[ 5]  1.00-2.00    sec    129 KBytes    1.05 Mbits/sec    0.390 ms      0/91 (0%)
[ 5]  2.00-3.00    sec    127 KBytes    1.04 Mbits/sec    0.255 ms      0/90 (0%)
[ 5]  3.00-4.00    sec    129 KBytes    1.05 Mbits/sec    0.163 ms      0/91 (0%)
[ 5]  4.00-5.00    sec    127 KBytes    1.04 Mbits/sec    0.153 ms      0/90 (0%)
[ 5]  5.00-6.00    sec    129 KBytes    1.05 Mbits/sec    0.172 ms      0/91 (0%)
[ 5]  6.00-7.00    sec    127 KBytes    1.04 Mbits/sec    0.273 ms      0/90 (0%)
[ 5]  7.00-8.00    sec    129 KBytes    1.05 Mbits/sec    0.347 ms      0/91 (0%)
[ 5]  8.00-9.00    sec    127 KBytes    1.04 Mbits/sec    0.265 ms      0/90 (0%)
[ 5]  9.00-10.00   sec    129 KBytes    1.05 Mbits/sec    0.183 ms      0/91 (0%)
[ 5] 10.00-10.04   sec     0.00 Bytes     0.00 bits/sec    0.183 ms      0/0 (0%)
-----
[ ID] Interval           Transfer     Bandwidth       Jitter        Lost/Total Datagrams
[ 5]  0.00-10.04   sec     0.00 Bytes     0.00 bits/sec    0.183 ms      0/897 (0%)
-----
Server listening on 5201
-----
```

The maximum throughput achieved is 1.04 Mbits/sec.

From the top command we noticed that CPU usage is not consumed and CPU is idle and I/O is also being affected. So by elimination, we can think that it is the memory.

We can also deduce that, both the VMs will have virtual memory from host and will have memory(RAM) crunch at the time of transferring packets. So, we think that that memory could be the bottleneck.

PROBLEM 4:

(Note: All the code and output is store in q4 folder submitted in zip)

README

Prerequisite: Install ansible on host machine

```
$ sudo apt-add-repository ppa:ansible/ansible
$ sudo apt-get update
$ sudo apt-get install ansible
$ ansible --version
ansible 2.7.0
```

1. Ansible Playbook: q4_1.yml

- a. Create pure L2 network and named khchoksi-netl2

```
ece792@ece792-Standard-PC-i440FX-PIIX-1996:~/linux_netw_hw/hw2$ virsh net-list --all
Name                               State    Autostart Persistent
-----
default                            active   yes        yes
khchoksi-netl2                     active   no         yes
khchoksiNETWORK2                   active   no         yes
khchoksiNETWORK3                   active   no         yes
```

- b. Create two VMs (khchoksi4, khchoksi5) and connect with this L2 Network
Please make sure to run the ansible script with X11 forwarding.
- The script will wait for configuration of VM using GUI (which can't be automated)

```
ece792@ece792-Standard-PC-i440FX-PIIX-1996:~/linux_netw_hw/hw2$ virsh list --all
Id   Name                               State
-----
3    khchoksi                           running
5    khchoksi1lab2VM2                   running
16   khchoksi4                           running
17   khchoksi5                           running
-    cloned_vm                           shut off
-    khchoksi3                           shut off
```

```
ece792@ece792-Standard-PC-i440FX-PIIX-1996:~/linux_netw_hw/hw2$ virsh domiflist khchoksi4
Interface Type      Source      Model      MAC
-----
vnet6     bridge   khchoksi-netl2 virtio     52:54:00:a2:4d:ea
vnet7     network  khchoksiNETWORK3 virtio     52:54:00:1c:7b:4c

ece792@ece792-Standard-PC-i440FX-PIIX-1996:~/linux_netw_hw/hw2$ virsh domiflist khchoksi5
Interface Type      Source      Model      MAC
-----
vnet8     bridge   khchoksi-netl2 virtio     52:54:00:21:ba:f1
vnet9     network  khchoksiNETWORK3 virtio     52:54:00:7c:f1:fe
```

Run playbook using following command

```
$ sudo ansible-playbook q4_1.yml
--extra-vars="/home/ece792/iso/CentOS-7-x86_64-Minimal-1804.iso"
```

Playbook: q4_1.yml:


```

---
- hosts: localhost
gather_facts: no
vars:
  network_name: khchoksi-net12
  bridge_name: sw2
  packages:
    - python-libvirt
    - python-lxml

  guests:
    - name: khchoksi4
      mem: 512
      vcpu: 1
      network: "{{ network_name }}"
    - name: khchoksi5
      mem: 1024
      vcpu: 2
      network: "{{ network_name }}"

  vm_disk_location: /var/lib/libvirt/images/
  iso_file_path: {{ iso_file | default('/home/ece792/iso/CentOS-7-x86_64-Minimal-1804.iso') }}

tasks:
  # Install required packages
  - name: Install required packages for libvirt, lxml
    apt:
      name: "{{ packages }}"
      become: yes

  # Define a new network
  - name: Define Virtual Network
    virt_net:
      command: define
      name: '{{ network_name }}'
      xml: '{{ lookup("template", "templates/bridge_template.xml.j2") }}'

  # Create and start a network
  - name: Create Virtual Network if not created
    virt_net:
      command: create
      name: "{{ network_name }}"
      ignore_errors: true

  # Stop a network
  # - name: Stop Virtual Network if running
  #   virt_net:
  #     command: stop
  #     name: "{{ network_name }}"
  #   ignore_errors: true # To make task idempotent

  # List available networks
  - name: List available networks
    virt_net:
      command: list_nets

```

```
# Create New VM and will pop up UI
- name: Create VM instance
  command: >
      virt-install -n {{ item.name }} -r {{ item.mem }} --vcpu={{ item.vcpu }} --cpu host --disk
path={{ vm_disk_location }}{{ item.name }}.img,size=5 --network network={{ network_name }} -c {{
iso_file_path }} -v
  become: yes
  with_items: "{{ guests }}"
```

Intermediate steps to setup ssh:

- I. Create a new NAT bridge virbr1 so that both the newly created vms can have ips.
- II. Create new network: khchoksiNETWORK3.xml

```
<network>
  <name>khchoksiNETWORK3</name>
  <uuid>eadcd6b7-c89a-43b5-9fe0-407eb0034038</uuid>
  <forward mode='nat' />
  <bridge name='virbr1' stp='on' delay='0' />
  <mac address='52:54:00:9f:f8:b6' />
  <ip address='192.168.119.1' netmask='255.255.255.0'>
    <dhcp>
      <range start='192.168.119.2' end='192.168.119.254' />
    </dhcp>
  </ip>
</network>
```

- III. \$ brctl addbr virbr1
 \$ virsh net- start khchoksiNETWORK3
 Add this interface to both the VMs and restart them.

```
<interface type='network'>
  <source network='khchoksiNETWORK3' />
  <model type='virtio' />
  <address type='pci' domain='0x0000' bus='0x00' slot='0x0a' function='0x0' />
</interface>
```

- IV. Get the ips of both the vms (if not assigned, do dhcpclient).
- V. Create SSH key pairs on host machine
 \$ ssh-keygen
 and then follow the command
- VI. Copy public keys to both guest machines
 \$ ssh-copy-id -i ./keys/vm_rsa.pub root@192.168.119.58
 \$ ssh-copy-id -i ./keys/vm_rsa.pub root@192.168.119.252
- VII. Create new inventory file as follows:

```
[vms]
localhost ansible_connection=local
192.168.119.58 ansible_ssh_user=root ansible_ssh_private_key_file=./keys/vm_rsa
192.168.119.252 ansible_ssh_user=root ansible_ssh_private_key_file=./keys/vm_rsa
```

2. Ansible playbook to collect logs:

- Make sure inventory file and keys are created as mentioned above.
- Run the playbook with total time(in minutes) as parameter given below. If time not defined, it will take default 5 minutes. Granularity is set as 1 minutes as mentioned in the description.

```
$ sudo ansible-playbook q4_2.yml -i ./inventory --extra-vars "time=7"
```

Logs will be generated at: /var/customlogs/logs.csv (Attached within zip:-> q4_2_logs.csv)

Note: Ansible playbook will create 'customlogs' directory if not present.

Playbook: q4_2.yml

```
---
- hosts: vms
  gather_facts: no
  vars:
    total_time: "{{ time | default(5) }}" # Defined total time if not passed from command line
    granularity: 60 #in seconds
    log_file_directory: /var/customlogs
    log_file_path: "{{ log_file_directory }}/logs.csv"
  tasks:
    - name: Create logs directory if not present
      file:
        path: "{{ log_file_directory }}"
        state: directory
        mode: 0777
        owner: ece792
        group: ece792
      become: yes
      delegate_to: localhost
      run_once: true

    - name: Generate Log CSV File Header
      shell: echo "hostname, timestamp, cpu1min, cpu5min, cpu15min" >> "{{ log_file_path }}"
      delegate_to: localhost
      run_once: true
      become: yes

    - name: Generate loop sequence based on input total time parameter
      set_fact:
        loop_sequence: "{{ loop_sequence | default([]) + [item | int] }}"
      with_sequence: start=1 end={{ total_time }}

# - name: debug_list
#   debug:
```

```

#         msg: "{{ loop_sequence }}"

- name: Fetch cpu usages from host and guests, store it in output variable
  shell: "echo -n '{{ hostvars[inventory_hostname]['inventory_hostname'] }},' && date +%X | awk -F,
'{{ printf \"%s, \", $1 }}' && uptime | sed 's/.*/load average: //' | awk -F\\, '{{ printf \"%s, %s, %s\\\", $1,
$2, $3 }}'"
  register: output
  loop: "{{ loop_sequence }}"
  loop_control:
    pause: "{{ granularity }}"

# - name: debugging
#   debug:
#     msg: "{{ item.stdout }}"
#   with_items: "{{ output.results }}"

- name: Writing logs to csv file
  shell: |
    echo "{{ item.stdout }}" >> "{{ log_file_path }}"
  with_items: "{{ output.results }}"
  delegate_to: localhost
  become: yes

```

Problem 5:

(Note: All the code and output is store in q5 folder submitted in zip)

1) Obtaining host information:

```
import sys
import libvirt
import random

conn = libvirt.open('qemu:///system')
if not conn:
    print "Connection failed"
    exit(1)

domainIDs = conn.listDomainsID()
if len(domainIDs) == 0:
    print "No active domains"
randomid = random.sample(domainIDs, 1)

rdom = conn.lookupByID(randomid[0])

state, maxmem, mem, cpus, cput = rdom.info()
print "UUID of the guest vm : ", rdom.UUIDString()
print "OS type of the guest vm : ", rdom.OSType()
print "Max vcpus of the guest vm : ", str(rdom.maxVcpus())
print "State of the guest vm : ", str(state)
print "Name of the guest vm : ", rdom.name()
print "Max memory of the guest vm : ", str(maxmem)
print "Number of cpus in the guest vm : ", str(cpus)

conn.close()
exit(1)
```

```
ece792@ece792-Standard-PC-i440FX-PIIX-1996:~/Downloads$ python host_info.py
Hostname : ece792-Standard-PC-i440FX-PIIX-1996
Number of vcpus : 16
Memory size : 24109
Clock speed of CPUs : 2199
Number of CPUs : 4
Virtualization type: QEMU
Canonical URI : qemu:///system
ece792@ece792-Standard-PC-i440FX-PIIX-1996:~/Downloads$
```

(README) To run this code, simply run following:

```
$ python host_info.py
```


2) Obtaining Guest Information

```
import sys
import libvirt

conn = libvirt.open('qemu:///system')

if conn == None:
    print('Connection failed')
    exit(1)

node_info = conn.getInfo()

print "Hostname : ", conn.getHostname()
print "Number of vcpus : ", conn.getMaxVcpus(None)
print "Memory size : ", node_info[1]
print "Clock speed of CPUs : ", node_info[3]
print "Number of CPUs : ", node_info[2]
print "Virtualization type: ", conn.getType()
print "Canonical URI : ", conn.getURI()

conn.close()
exit(1)
```

```
ece792@ece792-Standard-PC-i440FX-PIIX-1996:~/Downloads$ python guest_info.py
UUID of the guest vm : eda00eb7-7fe8-4360-8787-5c8bd76fe2d8
OS type of the guest vm : hvm
Max vcpus of the guest vm : 1
State of the guest vm : 1
Name of the guest vm : khchoksi4
Max memory of the guest vm : 524288
Number of cpus in the guest vm : 1
ece792@ece792-Standard-PC-i440FX-PIIX-1996:~/Downloads$
```

(README) To run this code, simply run following:

```
$ python guest_info.py
```

3) Performance Monitoring:

We calculate cpu utilization using : $100 * (\text{cpu_time}_{t_2} - \text{cpu_time}_{t_1}) / 10^9 * (t_2 - t_1)$

Where t1 and t2 are in seconds, and cpu_time is in nanoseconds

In cases where we have more than 1 vcpu, we divide the utilization by the number of vcpus to obtain the aggregate CPU utilization.

For memory, we calculate the utilization as : $1 - \text{memory_available} / \text{memory_actual}$

This gives us the amount of memory used, which we believe is a measure of the memory utilization

```
import sys
import libvirt
import random
import os
import datetime
import argparse
import collections
from time import sleep

py_parser = argparse.ArgumentParser(description='Monitor script')

py_parser.add_argument('order', nargs=1, choices = ["CPU", "MEM"], help="order to sort
by")
py_parser.add_argument('--threshold', nargs=1, type = float, help="threshold CPU
value")
py_parser.add_argument('--polling_interval', nargs=1, type = float, help="polling
interval value")
py_parser.add_argument('--moving_window', nargs=1, type = int, help="moving window
value")

py_args = py_parser.parse_args()

order = py_args.order[0]

if not py_args.threshold:
    threshold = 0
else:
    threshold = py_args.threshold[0]

conn = libvirt.open('qemu:///system')
if not conn:
    print "Connection failed"
    exit(1)

domainIDs = conn.listDomainsID()
if len(domainIDs) == 0:
    print "No active domains"

vm_id_list = [] + domainIDs
stats = []
for vm_id in vm_id_list:
```

```

    vm = conn.lookupByID(vm_id)
    cpu_stats = vm.getCPUStats(True)[0]
    mem_stats = vm.memoryStats()
    vcpus = vm.maxVcpus()
    stats.append([vm_id, cpu_stats['cpu_time']*1.0,\
1 - mem_stats['available']*1.0/mem_stats['actual'], vcpus, vm])

sleep(1)

for indx, stat in enumerate(stats):
    vm = stat[-1]
    cpu_stats = vm.getCPUStats(True)[0]
    stats[indx][1] = (cpu_stats['cpu_time'] - stats[indx][1])/10**9
    stats[indx][1] = (stats[indx][1]*100)/stats[indx][3]
    if stats[indx][1] > 100:
        stats[indx][1] = 100

#Sort by CPU or MEM
if order == "CPU":
    stats.sort(key=lambda x: x[1])
else:
    stats.sort(key=lambda x: x[2])

#Logging
if not os.path.isfile("alerts.csv"):
    log_file = open("alerts.csv", 'w')
    log_file.write("VM name, timestamp, CPU usage\n")
else:
    log_file = open("alerts.csv", 'a')

log_op = ""
print_op = ""
for vm_stat in stats:
    #Printing sorted list
    print "ID : ", vm_stat[0], " CPU usage : ", vm_stat[1]\
, " MEM usage : ", vm_stat[2]
    #If cpu > threshold, log and print
    if vm_stat[1] > threshold:
        log_op += vm_stat[-1].name()+"", "+str(datetime.datetime.now())+",
"+str(vm_stat[1])+"\n"
        print_op += vm_stat[-1].name()+"", "+str(datetime.datetime.now())+",
"+str(vm_stat[1])+"\n"

log_file.write(log_op)
print "\n", print_op

#Bonus part

```

```

poll_int = py_args.polling_interval[0]
mov_wind = py_args.moving_window[0]

if poll_int==None or mov_wind==None:
    exit(1)
poll_int = py_args.polling_interval[0]
mov_wind = py_args.moving_window[0]

prev_poll_time = {}
curr_poll_time = {}
polled_values = collections.defaultdict(list)

for indx, stat in enumerate(stats):
    vm = stat[-1]
    cpu_stats = vm.getCPUStats(True)[0]
    prev_poll_time[stat[0]] = cpu_stats['cpu_time']

if not os.path.isfile("mov_avgs.csv"):
    mavgs = open("mov_avgs.csv", 'w')
    mavgs.write("VM ID, Moving average CPU usage\n")
else:
    mavgs = open("mov_avgs.csv", 'a')

poll_timer = 0
try:
    while True:
        sleep(poll_int)

        if poll_timer >= mov_wind*poll_int:
            #Log the moving window averages

            unsorted_list = []
            for v in polled_values:
                unsorted_list.append([v,
sum(polled_values[v])/len(polled_values[v])])
                polled_values[v].pop(0)
            unsorted_list.sort(key = lambda x: x[1])
            for v in unsorted_list:
                mavgs.write(str(v[0])+", "+str(v[1])+"\n")
            #print unsorted_list

            for indx, stat in enumerate(stats):
                vm = stat[-1]
                cpu_stats = vm.getCPUStats(True)[0]
                curr_poll_time[stat[0]] = cpu_stats['cpu_time']
                polled_values[stat[0]].append(100*(curr_poll_time[stat[0]] - \
prev_poll_time[stat[0]])/(10**9 * poll_int * stats[indx][3]))

```

```
        prev_poll_time[stat[0]] = curr_poll_time[stat[0]]
        #print polled_values
        poll_timer += poll_int
except:
    mavgs.close()
    exit(0)
```

```
ece792@ece792-Standard-PC-i440FX-PIIX-1996:~/Downloads$ python mon
itor.py CPU --threshold 0.6
ID :   23   CPU usage :  0.2307952   MEM usage :  0.102426528931
ID :   27   CPU usage :  0.924702175   MEM usage :  0.102426528931
ID :   26   CPU usage :  1.77755225   MEM usage :  0.0317306518555
ID :   25   CPU usage :  2.6463952   MEM usage :  0.0475387573242

khchoksilab2VM2, 2018-10-16 22:14:15.122545, 0.924702175
khchoksi5, 2018-10-16 22:14:15.122639, 1.77755225
khchoksi4, 2018-10-16 22:14:15.122722, 2.6463952
```



```
ece792@ece792-Standard-PC-i440FX-PIIX-1996:~/Downloads$ python monitor.py CPU --threshold 0.6 --polling_interval 1 --moving_window 3
```

```
ID : 23 CPU usage : 0.1499394 MEM usage : 0.102426528931
ID : 26 CPU usage : 0.3936715 MEM usage : 0.0317306518555
ID : 27 CPU usage : 0.878282775 MEM usage : 0.102426528931
ID : 25 CPU usage : 0.94446 MEM usage : 0.0475387573242
```

```
khchoksilab2VM2, 2018-10-16 22:19:54.724152, 0.878282775
khchoksi4, 2018-10-16 22:19:54.724198, 0.94446
```

```
ece792@ece792-Standard-PC-i440FX-PIIX-1996:~/Downloads$ cat alerts.csv
```

```
VM name, timestamp, CPU usage
```

```
khchoksilab2VM2, 2018-10-16 22:19:54.724073, 0.878282775
```

```
khchoksi4, 2018-10-16 22:19:54.724183, 0.94446
```

```
ece792@ece792-Standard-PC-i440FX-PIIX-1996:~/Downloads$ cat mavg
```

```
cat: mavg: No such file or directory
```

```
ece792@ece792-Standard-PC-i440FX-PIIX-1996:~/Downloads$ cat mov_avgs.csv
```

```
VM name, Moving average CPU usage
```

```
26, 0.465346466667
```

```
27, 0.619114116667
```

```
25, 0.7367698
```

```
23, 2.02397663333
```

```
26, 0.338603416667
```

```
23, 0.55610685
```

```
25, 0.644380266667
```

```
27, 0.684600033333
```

```
26, 0.452477316667
```

```
23, 0.558218691667
```

```
27, 0.604018258333
```

```
25, 0.8326495
```

```
23, 0.210707616667
```

```
26, 0.4832872
```

```
25, 0.637493
```


```
27, 0.645893933333
```

```
23, 0.2299261
```

```
26, 0.508174866667
```

```
27, 0.6301075
```

PROBLEM 6:



Distributed vSwitch: Distributed vSwitches allow different hosts to use the switch as long as they exist within the same host cluster. A distributed vSwitch extends its ports and management across all the servers in a cluster, supporting up to 500 hosts per distributed switch. Instead of making virtual networks more complicated with its additional options, the distributed vSwitch simplifies operations and helps catch configuration errors and increase network visibility.

Standard vSwitch: A standard vSwitch works within one ESX/ESXi host only. Standard switch is created in host level i.e. we can create and manage vSphere standard switch independently on an ESXi host. Inbound traffic shaping is not available as a part in the standard switch.

In cases where we want to extend the L2 layer across VMs in different hosts, instead of creating GRE/VXLAN tunnels between each host pair, we could use the distributed vSwitch.

Ref:
<https://searchvmware.techtarget.com/photostory/2240185944/Getting-VMware-terminology-straight/9/How-do-switches-vSwitches-and-distributed-vSwitches-differ>

PROBLEM 7:

1. If 2 VMs connected to same bridge in bridge mode:

a. Same MAC address:

It will overwrite the entries for same MAC address when ARP will be done. The 2VMs should be able to ping each other.

As they have different IP address, the ARP will be resolved and ping will be successful between two VMs as well as bridge will work fine.

```
ens9: flags=4163<UP,BROADCAST,RUNNING,MULTICAST> mtu 1500
    inet 10.0.0.2 netmask 255.255.255.0 broadcast 10.0.0.255
    ether 52:54:00:ab:7f:fb txqueuelen 1000 (Ethernet)
    RX packets 37122 bytes 52521987 (50.0 MiB)
    RX errors 0 dropped 0 overruns 0 frame 0
    TX packets 13363 bytes 946315 (924.1 KiB)
    TX errors 0 dropped 0 overruns 0 carrier 0 collisions 0

eth0: flags=4163<UP,BROADCAST,RUNNING,MULTICAST> mtu 1500
    inet 192.168.122.92 netmask 255.255.255.0 broadcast 192.168.122.255
    ether 52:54:00:2e:29:b6 txqueuelen 1000 (Ethernet)
    RX packets 187169 bytes 10760135 (10.2 MiB)
    RX errors 0 dropped 20 overruns 0 frame 0
    TX packets 18472 bytes 1863201 (1.7 MiB)
    TX errors 0 dropped 0 overruns 0 carrier 0 collisions 0

eth1: flags=4163<UP,BROADCAST,RUNNING,MULTICAST> mtu 1500
    inet 10.0.0.5 netmask 255.255.255.0 broadcast 10.0.0.255
    ether 52:54:00:dd:70:cb txqueuelen 1000 (Ethernet)
    RX packets 2533 bytes 164265 (160.4 KiB)
    RX errors 0 dropped 0 overruns 0 frame 0
    TX packets 0 bytes 0 (0.0 B)
    TX errors 0 dropped 0 overruns 0 carrier 0 collisions 0

lo: flags=73<UP,LOOPBACK,RUNNING> mtu 65536
    inet 127.0.0.1 netmask 255.0.0.0
    inet6 ::1 prefixlen 128 scopeid 0x10<host>
    loop txqueuelen 1000 (Local Loopback)
    RX packets 563 bytes 63118 (61.6 KiB)
    RX errors 0 dropped 0 overruns 0 frame 0
    TX packets 563 bytes 63118 (61.6 KiB)
    TX errors 0 dropped 0 overruns 0 carrier 0 collisions 0

eth0: flags=4163<UP,BROADCAST,RUNNING,MULTICAST> mtu 1500
    inet 192.168.122.51 netmask 255.255.255.0 broadcast 192.168.122.255
    ether 52:54:00:cd:dc:14 txqueuelen 1000 (Ethernet)
    RX packets 903 bytes 70228 (68.5 KiB)
    RX errors 0 dropped 7 overruns 0 frame 0
    TX packets 494 bytes 64107 (62.6 KiB)
    TX errors 0 dropped 0 overruns 0 carrier 0 collisions 0

eth1: flags=4163<UP,BROADCAST,RUNNING,MULTICAST> mtu 1500
    inet 10.0.0.3 netmask 255.255.255.0 broadcast 10.0.0.255
    ether 52:54:00:dd:70:cb txqueuelen 1000 (Ethernet)
    RX packets 421 bytes 25307 (24.7 KiB)
    RX errors 0 dropped 7 overruns 0 frame 0
    TX packets 163 bytes 21758 (21.2 KiB)
    TX errors 0 dropped 0 overruns 0 carrier 0 collisions 0

lo: flags=73<UP,LOOPBACK,RUNNING> mtu 65536
    inet 127.0.0.1 netmask 255.0.0.0
    inet6 ::1 prefixlen 128 scopeid 0x10<host>
    loop txqueuelen 1000 (Local Loopback)
    RX packets 0 bytes 0 (0.0 B)
    RX errors 0 dropped 0 overruns 0 frame 0
    TX packets 0 bytes 0 (0.0 B)
    TX errors 0 dropped 0 overruns 0 carrier 0 collisions 0

[root@localhost ~]# ping 10.0.0.5 -c 2
PING 10.0.0.5 (10.0.0.5) 56(84) bytes of data.
64 bytes from 10.0.0.5: icmp_seq=1 ttl=64 time=1.38 ms
64 bytes from 10.0.0.5: icmp_seq=2 ttl=64 time=1.84 ms

--- 10.0.0.5 ping statistics ---
 2 packets transmitted, 2 received, 0% packet loss, time 1001ms
rtt min/avg/max/mdev = 1.385/1.613/1.842/0.232 ms
```

b. Same IP address:

Switch will throw an error that, duplicate use of ip detected.

```
245 1096.123639937 RealtekU_dd:70:cb → Broadcast ARP 42 Who has 10.0.0.5? Tell 10.0.0.2 (duplicate use of 10.0.0.2 detected!)
)
246 1097.125023406 RealtekU_dd:70:cb → Broadcast ARP 42 Who has 10.0.0.5? Tell 10.0.0.2 (duplicate use of 10.0.0.2 detected!)
)
247 1098.127041832 RealtekU_dd:70:cb → Broadcast ARP 42 Who has 10.0.0.5? Tell 10.0.0.2 (duplicate use of 10.0.0.2 detected!)
)
248 1100.124554352 RealtekU_dd:70:cb → Broadcast ARP 42 Who has 10.0.0.5? Tell 10.0.0.2 (duplicate use of 10.0.0.2 detected!)
)
249 1101.127078045 RealtekU_dd:70:cb → Broadcast ARP 42 Who has 10.0.0.5? Tell 10.0.0.2 (duplicate use of 10.0.0.2 detected!)
)
250 1102.129136591 RealtekU_dd:70:cb → Broadcast ARP 42 Who has 10.0.0.5? Tell 10.0.0.2 (duplicate use of 10.0.0.2 detected!)
)
251 1104.125752241 RealtekU_dd:70:cb → Broadcast ARP 42 Who has 10.0.0.5? Tell 10.0.0.2 (duplicate use of 10.0.0.2 detected!)
```

```

ens9: flags=4163<UP,BROADCAST,RUNNING,MULTICAST> mtu 1500
    inet 10.0.0.2 netmask 255.255.255.0 broadcast 10.0.0.255
    ether 52:54:00:ab:7f:fb txqueuelen 1000 (Ethernet)
    RX packets 37142 bytes 52524903 (50.0 MiB)
    RX errors 0 dropped 0 overruns 0 frame 0
    TX packets 13383 bytes 947603 (925.3 KiB)
    TX errors 0 dropped 0 overruns 0 carrier 0 collisions 0

eth0: flags=4163<UP,BROADCAST,RUNNING,MULTICAST> mtu 1500
    inet 192.168.122.92 netmask 255.255.255.0 broadcast 192.168.122.255
    ether 52:54:00:2e:29:b6 txqueuelen 1000 (Ethernet)
    RX packets 188193 bytes 10832971 (10.3 MiB)
    RX errors 0 dropped 20 overruns 0 frame 0
    TX packets 18844 bytes 1909021 (1.8 MiB)
    TX errors 0 dropped 0 overruns 0 carrier 0 collisions 0

eth1: flags=4163<UP,BROADCAST,RUNNING,MULTICAST> mtu 1500
    inet 10.0.0.5 netmask 255.255.255.0 broadcast 10.0.0.255
    ether 52:54:00:dd:70:cb txqueuelen 1000 (Ethernet)
    RX packets 2979 bytes 189155 (184.7 KiB)
    RX errors 0 dropped 0 overruns 0 frame 0
    TX packets 0 bytes 0 (0.0 B)
    TX errors 0 dropped 0 overruns 0 carrier 0 collisions 0

lo: flags=73<UP,LOOPBACK,RUNNING> mtu 65536
    inet 127.0.0.1 netmask 255.0.0.0
    inet6 ::1 prefixlen 128 scopeid 0x10<host>
    loop txqueuelen 1000 (Local Loopback)
    RX packets 622 bytes 68718 (67.1 KiB)
    RX errors 0 dropped 0 overruns 0 frame 0
    TX packets 622 bytes 68718 (67.1 KiB)
    TX errors 0 dropped 0 overruns 0 carrier 0 collisions 0

[root@localhost ~]# ping 10.0.0.5 -c 2
PING 10.0.0.5 (10.0.0.5) 56(84) bytes of data.
From 10.0.0.2 icmp_seq=1 Destination Host Unreachable
From 10.0.0.2 icmp_seq=2 Destination Host Unreachable

--- 10.0.0.5 ping statistics ---
 2 packets transmitted, 0 received, +2 errors, 100% packet loss, time
999ms
pipe 2

```


2. If 2 Vms connected to different bridge (both bridge mode)

a. Same MAC address:

As both VMs are in different networks, it won't affect the bridge.

```
[root@localhost ~]# ifconfig
eth0: flags=4163<UP,BROADCAST,RUNNING,MULTICAST> mtu 1500
    inet 192.168.122.51 netmask 255.255.255.0 broadcast 192.168.122.255
    ether 52:54:00:cd:dc:14 txqueuelen 1000 (Ethernet)
    RX packets 712 bytes 55417 (54.1 KiB)
    RX errors 0 dropped 9 overruns 0 frame 0
    TX packets 291 bytes 35399 (34.5 KiB)
    TX errors 0 dropped 0 overruns 0 carrier 0 collisions 0

eth1: flags=4163<UP,BROADCAST,RUNNING,MULTICAST> mtu 1500
    inet 10.0.0.3 netmask 255.255.255.0 broadcast 10.0.0.255
    inet6 fe80::fe25:87bb:ac80:78fe prefixlen 64 scopeid 0x2
link> ether 52:54:00:21:ba:f1 txqueuelen 1000 (Ethernet)
    RX packets 279 bytes 14872 (14.5 KiB)
    RX errors 0 dropped 8 overruns 0 frame 0
    TX packets 94 bytes 16392 (16.0 KiB)
    TX errors 0 dropped 0 overruns 0 carrier 0 collisions 0

lo: flags=73<UP,LOOPBACK,RUNNING> mtu 65536
    inet 127.0.0.1 netmask 255.0.0.0
    inet6 ::1 prefixlen 128 scopeid 0x10<host>
    loop txqueuelen 1000 (Local Loopback)
    RX packets 4 bytes 448 (448.0 B)
    RX errors 0 dropped 0 overruns 0 frame 0
    TX packets 4 bytes 448 (448.0 B)
    TX errors 0 dropped 0 overruns 0 carrier 0 collisions 0

[root@localhost ~]# ping 10.0.0.5
PING 10.0.0.5 (10.0.0.5) 56(84) bytes of data.
64 bytes from 10.0.0.5: icmp_seq=1 ttl=64 time=0.858 ms
64 bytes from 10.0.0.5: icmp_seq=2 ttl=64 time=2.70 ms
64 bytes from 10.0.0.5: icmp_seq=3 ttl=64 time=3.06 ms
^C
--- 10.0.0.5 ping statistics ---
3 packets transmitted, 3 received, 0% packet loss, time 2003ms
rtt min/avg/max/mdev = 0.858/2.206/3.061/0.965 ms
[root@localhost ~]#
```

```
[root@localhost ~]# ifconfig
eth0: flags=4163<UP,BROADCAST,RUNNING,MULTICAST> mtu 1500
    inet 12.0.0.3 netmask 255.255.255.0 broadcast 12.0.0.255
    ether 52:54:00:21:ba:f1 txqueuelen 1000 (Ethernet)
    RX packets 46 bytes 4016 (3.9 KiB)
    RX errors 0 dropped 0 overruns 0 frame 0
    TX packets 5 bytes 434 (434.0 B)
    TX errors 0 dropped 0 overruns 0 carrier 0 collisions 0

eth1: flags=4163<UP,BROADCAST,RUNNING,MULTICAST> mtu 1500
    inet 192.168.119.252 netmask 255.255.255.0 broadcast 192.168.119.255
    inet6 fe80::b572:f8d4:deb0:e754 prefixlen 64 scopeid 0x20<
link> ether 52:54:00:7c:f1:fe txqueuelen 1000 (Ethernet)
    RX packets 63378 bytes 22656519 (21.6 MiB)
    RX errors 0 dropped 20 overruns 0 frame 0
    TX packets 25688 bytes 2230714 (2.1 MiB)
    TX errors 0 dropped 0 overruns 0 carrier 0 collisions 0

lo: flags=73<UP,LOOPBACK,RUNNING> mtu 65536
    inet 127.0.0.1 netmask 255.0.0.0
    inet6 ::1 prefixlen 128 scopeid 0x10<host>
    loop txqueuelen 1000 (Local Loopback)
    RX packets 68 bytes 5920 (5.7 KiB)
    RX errors 0 dropped 0 overruns 0 frame 0
    TX packets 68 bytes 5920 (5.7 KiB)
    TX errors 0 dropped 0 overruns 0 carrier 0 collisions 0

[root@localhost ~]# ping 12.0.0.2
PING 12.0.0.2 (12.0.0.2) 56(84) bytes of data.
64 bytes from 12.0.0.2: icmp_seq=1 ttl=64 time=0.564 ms
64 bytes from 12.0.0.2: icmp_seq=2 ttl=64 time=0.603 ms
64 bytes from 12.0.0.2: icmp_seq=3 ttl=64 time=0.566 ms
^C
--- 12.0.0.2 ping statistics ---
3 packets transmitted, 3 received, 0% packet loss, time 2001ms
rtt min/avg/max/mdev = 0.564/0.577/0.603/0.033 ms
[root@localhost ~]#
```

b. Same IP address:

As both the VMs are in different network, it won't allow to change the ip of second VM and network won't be affected.

```
    inet 192.168.122.51 netmask 255.255.255.0 broadcast 192.168.122.255
    ether 52:54:00:cd:dc:14 txqueuelen 1000 (Ethernet)
    RX packets 1071 bytes 81003 (79.1 KiB)
    RX errors 0 dropped 9 overruns 0 frame 0
    TX packets 433 bytes 53579 (52.3 KiB)
    TX errors 0 dropped 0 overruns 0 carrier 0 collisions 0

eth1: flags=4163<UP,BROADCAST,RUNNING,MULTICAST> mtu 1500
    inet 10.0.0.3 netmask 255.255.255.0 broadcast 10.0.0.255
    ether 52:54:00:21:ba:f1 txqueuelen 1000 (Ethernet)
    RX packets 416 bytes 22206 (21.6 KiB)
    RX errors 0 dropped 8 overruns 0 frame 0
    TX packets 139 bytes 23626 (23.0 KiB)
    TX errors 0 dropped 0 overruns 0 carrier 0 collisions 0

lo: flags=73<UP,LOOPBACK,RUNNING> mtu 65536
    inet 127.0.0.1 netmask 255.0.0.0
    inet6 ::1 prefixlen 128 scopeid 0x10<host>
    loop txqueuelen 1000 (Local Loopback)
    RX packets 4 bytes 448 (448.0 B)
    RX errors 0 dropped 0 overruns 0 frame 0
    TX packets 4 bytes 448 (448.0 B)
    TX errors 0 dropped 0 overruns 0 carrier 0 collisions 0

[root@localhost ~]# ping 10.0.0.5 -c 2
PING 10.0.0.5 (10.0.0.5) 56(84) bytes of data.
64 bytes from 10.0.0.5: icmp_seq=1 ttl=64 time=1.15 ms
64 bytes from 10.0.0.5: icmp_seq=2 ttl=64 time=2.04 ms

--- 10.0.0.5 ping statistics ---
2 packets transmitted, 2 received, 0% packet loss, time 1001ms
rtt min/avg/max/mdev = 1.150/1.597/2.044/0.447 ms
[root@localhost ~]#
```

```
eth0: flags=4163<UP,BROADCAST,RUNNING,MULTICAST> mtu 1500
    inet 10.0.0.3 netmask 255.255.255.0 broadcast 10.0.0.255
    ether 52:54:00:21:ba:f1 txqueuelen 1000 (Ethernet)
    RX packets 58 bytes 4912 (4.7 KiB)
    RX errors 0 dropped 0 overruns 0 frame 0
    TX packets 17 bytes 1330 (1.2 KiB)
    TX errors 0 dropped 0 overruns 0 carrier 0 collisions 0

eth1: flags=4163<UP,BROADCAST,RUNNING,MULTICAST> mtu 1500
    inet 192.168.119.252 netmask 255.255.255.0 broadcast 192.168.119.255
    inet6 fe80::b572:f8d4:deb0:e754 prefixlen 64 scopeid 0x20<
link> ether 52:54:00:7c:f1:fe txqueuelen 1000 (Ethernet)
    RX packets 64070 bytes 22710771 (21.6 MiB)
    RX errors 0 dropped 20 overruns 0 frame 0
    TX packets 26012 bytes 2271557 (2.1 MiB)
    TX errors 0 dropped 0 overruns 0 carrier 0 collisions 0

lo: flags=73<UP,LOOPBACK,RUNNING> mtu 65536
    inet 127.0.0.1 netmask 255.0.0.0
    inet6 ::1 prefixlen 128 scopeid 0x10<host>
    loop txqueuelen 1000 (Local Loopback)
    RX packets 68 bytes 5920 (5.7 KiB)
    RX errors 0 dropped 0 overruns 0 frame 0
    TX packets 68 bytes 5920 (5.7 KiB)
    TX errors 0 dropped 0 overruns 0 carrier 0 collisions 0

[root@localhost ~]# ping 10.0.0.10 -c 2
PING 10.0.0.10 (10.0.0.10) 56(84) bytes of data.
64 bytes from 10.0.0.10: icmp_seq=1 ttl=64 time=0.542 ms
64 bytes from 10.0.0.10: icmp_seq=2 ttl=64 time=0.798 ms

--- 10.0.0.10 ping statistics ---
2 packets transmitted, 2 received, 0% packet loss, time 1000ms
rtt min/avg/max/mdev = 0.542/0.670/0.798/0.128 ms
[root@localhost ~]#
```


3. If 2 Vms connected to different bridge (both routed mode)

- a. **Same MAC address:** If the two VMs are connected to different bridges (over here sw3 and sw4) then everything will work correctly even if we provide same MAC address to both the VMs. As they both are on different networks (routed_subnet1) and (routed_subnet2), it won't hinder each other.

```
ece792@ece792-Standard-PC-i440FX-PIIX-1996:~$ virsh domiflist khchoksilab2VM
Interface  Type      Source      Model      MAC
-----
vnet0      network   default     virtio      52:54:00:cd:dc:14
vnet1      network   routed_net1 virtio      52:54:00:84:18:43

ece792@ece792-Standard-PC-i440FX-PIIX-1996:~$ virsh domiflist khchoksi4
Interface  Type      Source      Model      MAC
-----
vnet2      bridge    khchoksi-netl2 virtio      52:54:00:a2:4d:ea
vnet3      network   khchoksiNETWORK3 virtio      52:54:00:1c:7b:4c
vnet6      network   routed_net2 virtio      52:54:00:84:18:43
```

```
ece792@ece792-Standard-PC-i440FX-PIIX-1996:~$ ssh root@192.168.122.51
root@192.168.122.51's password:
Last login: Tue Oct 16 21:55:01 2018
[root@localhost ~]# ifconfig
eth0: flags=4163<UP,BROADCAST,RUNNING,MULTICAST> mtu 1500
    inet 192.168.122.51 netmask 255.255.255.0 broadcast 192.168.122.255
    ether 52:54:00:cd:dc:14 txqueuelen 1000 (Ethernet)
    RX packets 158 bytes 13883 (13.5 KiB)
    RX errors 0 dropped 8 overruns 0 frame 0
    TX packets 36 bytes 5657 (5.5 KiB)
    TX errors 0 dropped 0 overruns 0 carrier 0 collisions 0

eth1: flags=4163<UP,BROADCAST,RUNNING,MULTICAST> mtu 1500
    inet 192.168.120.149 netmask 255.255.255.0 broadcast 192.168.120.255
    inet6 fe80::70c8:d1c:88e0:7408 prefixlen 64 scopeid 0x20<link>
    ether 52:54:00:84:18:43 txqueuelen 1000 (Ethernet)
    RX packets 125 bytes 11082 (10.8 KiB)
    RX errors 0 dropped 8 overruns 0 frame 0
    TX packets 62 bytes 6688 (6.5 KiB)
    TX errors 0 dropped 0 overruns 0 carrier 0 collisions 0

lo: flags=73<UP,LOOPBACK,RUNNING> mtu 65536
    inet 127.0.0.1 netmask 255.0.0.0
    inet6 ::1 prefixlen 128 scopeid 0x10<host>
    loop txqueuelen 1000 (Local Loopback)
    RX packets 0 bytes 0 (0.0 B)
    RX errors 0 dropped 0 overruns 0 frame 0
    TX packets 0 bytes 0 (0.0 B)
    TX errors 0 dropped 0 overruns 0 carrier 0 collisions 0

[root@localhost ~]#
[root@localhost ~]#
[root@localhost ~]#
[root@localhost ~]#
[root@localhost ~]#
[root@localhost ~]#
[root@localhost ~]#
[root@localhost ~]#
[root@localhost ~]#
```

```
ece792@ece792-Standard-PC-i440FX-PIIX-1996:~$ ssh root@192.168.119.58
root@192.168.119.58's password:
Last login: Tue Oct 16 19:57:21 2018 from gateway
[root@localhost ~]#
[root@localhost ~]#
[root@localhost ~]#
[root@localhost ~]# ifconfig
eth0: flags=4163<UP,BROADCAST,RUNNING,MULTICAST> mtu 1500
    ether 52:54:00:a2:4d:ea txqueuelen 1000 (Ethernet)
    RX packets 1 bytes 110 (110.0 B)
    RX errors 0 dropped 0 overruns 0 frame 0
    TX packets 0 bytes 0 (0.0 B)
    TX errors 0 dropped 0 overruns 0 carrier 0 collisions 0

eth1: flags=4163<UP,BROADCAST,RUNNING,MULTICAST> mtu 1500
    inet 192.168.119.58 netmask 255.255.255.0 broadcast 192.168.119.255
    inet6 fe80::b6c9:828e:bb70:314a prefixlen 64 scopeid 0x20<link>
    ether 52:54:00:1c:7b:4c txqueuelen 1000 (Ethernet)
    RX packets 238 bytes 21926 (21.4 KiB)
    RX errors 0 dropped 10 overruns 0 frame 0
    TX packets 124 bytes 13018 (12.7 KiB)
    TX errors 0 dropped 0 overruns 0 carrier 0 collisions 0

eth2: flags=4163<UP,BROADCAST,RUNNING,MULTICAST> mtu 1500
    inet 192.168.121.149 netmask 255.255.255.0 broadcast 192.168.121.255
    inet6 fe80::142b:12a8:dc1b:3dcf prefixlen 64 scopeid 0x20<link>
    ether 52:54:00:84:18:43 txqueuelen 1000 (Ethernet)
    RX packets 132 bytes 10087 (9.8 KiB)
    RX errors 0 dropped 10 overruns 0 frame 0
    TX packets 10 bytes 1288 (1.2 KiB)
    TX errors 0 dropped 0 overruns 0 carrier 0 collisions 0

lo: flags=73<UP,LOOPBACK,RUNNING> mtu 65536
    inet 127.0.0.1 netmask 255.0.0.0
    inet6 ::1 prefixlen 128 scopeid 0x10<host>
    loop txqueuelen 1000 (Local Loopback)
    RX packets 0 bytes 0 (0.0 B)
    RX errors 0 dropped 0 overruns 0 frame 0
    TX packets 0 bytes 0 (0.0 B)
    TX errors 0 dropped 0 overruns 0 carrier 0 collisions 0
```

b. Same IP address:

If they have the same IP address, the VM which has its IP address from the other subnet will lose connectivity with the subnet it is in. In our case, the VM khchoksilab2VM2 will lose connectivity from the 192.168.120.0/24 network as it has been moved to the 192.168.121.0/24 network IP address.

```

ece792@ece792-Standard-PC-i440FX-PIIX-1996:/etc/libvirt/qemu$ virsh
domiflist khchoksi

```

Interface	Type	Source	Model	MAC
vnet4	network	default	virtio	52:54:00:2e:29:b6
vnet5	network	routed_net1	virtio	52:54:00:dd:70:cb

```

ece792@ece792-Standard-PC-i440FX-PIIX-1996:/etc/libvirt/qemu$ virsh
domiflist khchoksilab2VM2

```

Interface	Type	Source	Model	MAC
vnet0	network	default	virtio	52:54:00:cd:dc:14
vnet1	network	routed_net1	virtio	52:54:00:84:18:43

```

ece792@ece792-Standard-PC-i440FX-PIIX-1996:/etc/libvirt/qemu$ virsh
domiflist khchoksi4

```

Interface	Type	Source	Model	MAC
vnet2	bridge	khchoksi-net12	virtio	52:54:00:a2:4d:ea
vnet3	network	khchoksiNETWORK3	virtio	52:54:00:1c:7b:4c
vnet6	network	routed_net2	virtio	52:54:00:84:18:43

```

ece792@ece792-Standard-PC-i440FX-PIIX-1996:/etc/libvirt/qemu$ virsh
domiflist khchoksi5

```

Interface	Type	Source	Model	MAC
vnet7	bridge	khchoksi-net12	virtio	52:54:00:21:ba:f1
vnet8	network	khchoksiNETWORK3	virtio	52:54:00:7c:f1:fe
vnet9	network	routed_net2	virtio	52:54:00:af:e1:6d

```

[root@localhost ~]# ifconfig
eth0: flags=4163<UP,BROADCAST,RUNNING,MULTICAST> mtu 1500
    inet 192.168.122.51 netmask 255.255.255.0 broadcast 192.168.122.255
    ether 52:54:00:cd:dc:14 txqueuelen 1000 (Ethernet)
    RX packets 478 bytes 32847 (32.0 KiB)
    RX errors 0 dropped 8 overruns 0 frame 0
    TX packets 70 bytes 11637 (11.3 KiB)
    TX errors 0 dropped 0 overruns 0 carrier 0 collisions 0

eth1: flags=4163<UP,BROADCAST,RUNNING,MULTICAST> mtu 1500
    inet 192.168.121.149 netmask 255.255.255.0 broadcast 192.168.121.255
    inet6 fe80::70c8:d11c:88e0:7408 prefixlen 64 scopeid 0x20<link>
    ether 52:54:00:84:18:43 txqueuelen 1000 (Ethernet)
    RX packets 385 bytes 24952 (24.3 KiB)
    RX errors 0 dropped 8 overruns 0 frame 0
    TX packets 74 bytes 7650 (7.4 KiB)
    TX errors 0 dropped 0 overruns 0 carrier 0 collisions 0

lo: flags=73<UP,LOOPBACK,RUNNING> mtu 65536
    inet 127.0.0.1 netmask 255.0.0.0
    inet6 ::1 prefixlen 128 scopeid 0x10<host>
    loop txqueuelen 1000 (Local Loopback)
    RX packets 0 bytes 0 (0.0 B)
    RX errors 0 dropped 0 overruns 0 frame 0
    TX packets 0 bytes 0 (0.0 B)

```